

# Progress of Pre-training in NLP: from word embedding to BERT

Jun Chen

2018/11/10

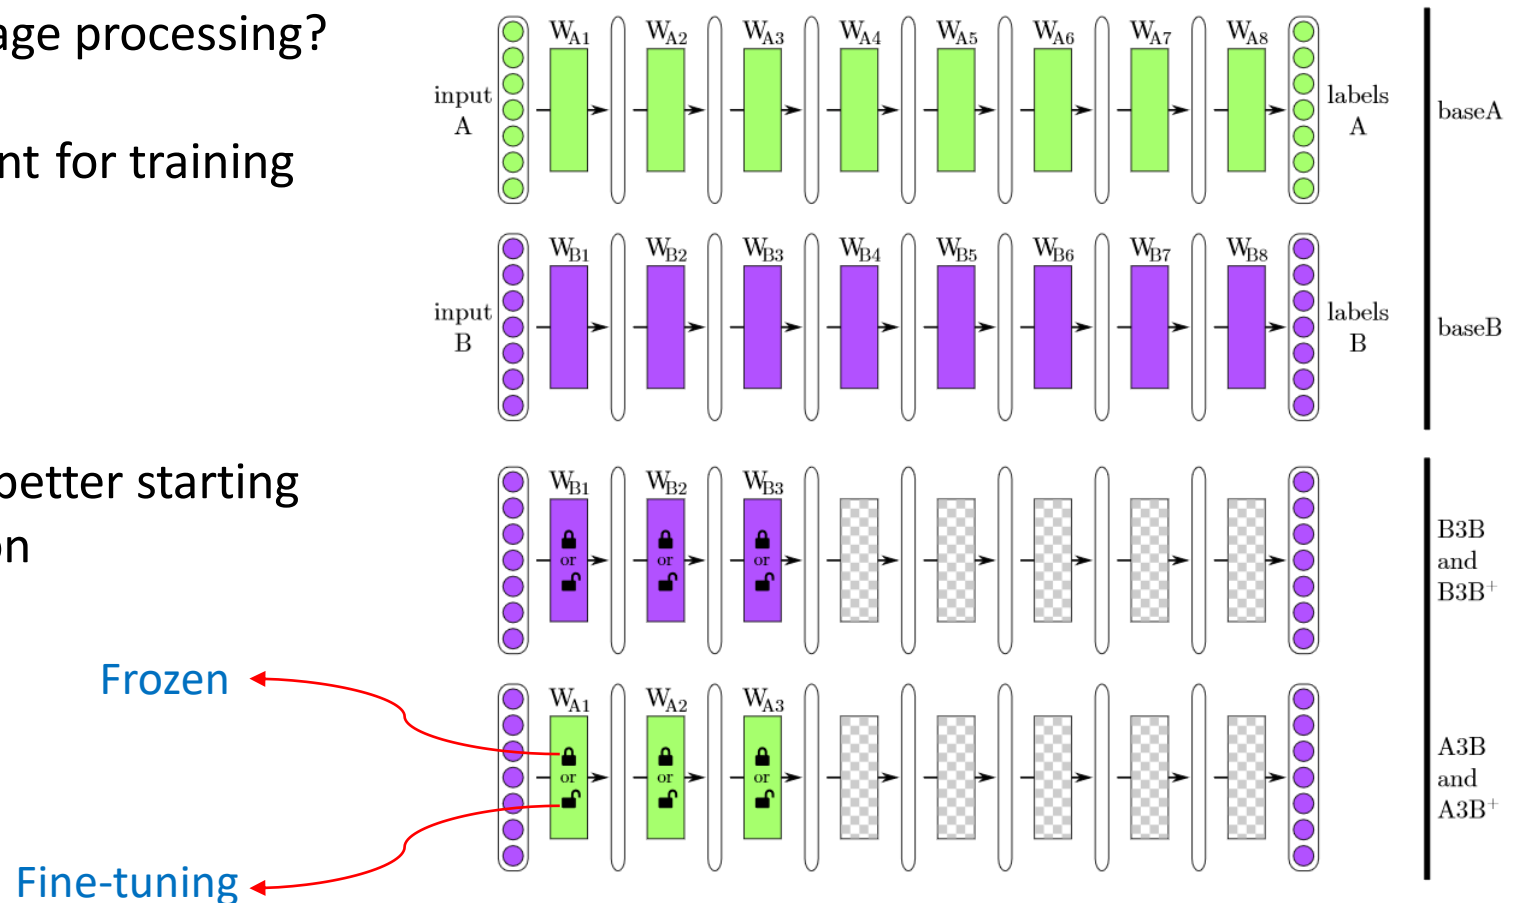
# Table of contents

- Pre-training in image processing
- Language model to word embedding
- Word embedding to ELMO
- Word embedding to GTP
- The birth of BERT

# Pre-training in image processing

Why pre-training is popular in image processing?

1. Small training set is not sufficient for training complex neural network
2. Speed up the training process
3. Parameter initialization: find a better starting point; be beneficial to optimization



# Pre-training in image processing

## Why freeze the weights?


First few layers capture universal features like curves and edges

Deep neural networks learn hierarchical feature representations



## Why fine-tuning?

Subsequent layers capture features more specific to task/dataset



# Pre-training in image processing

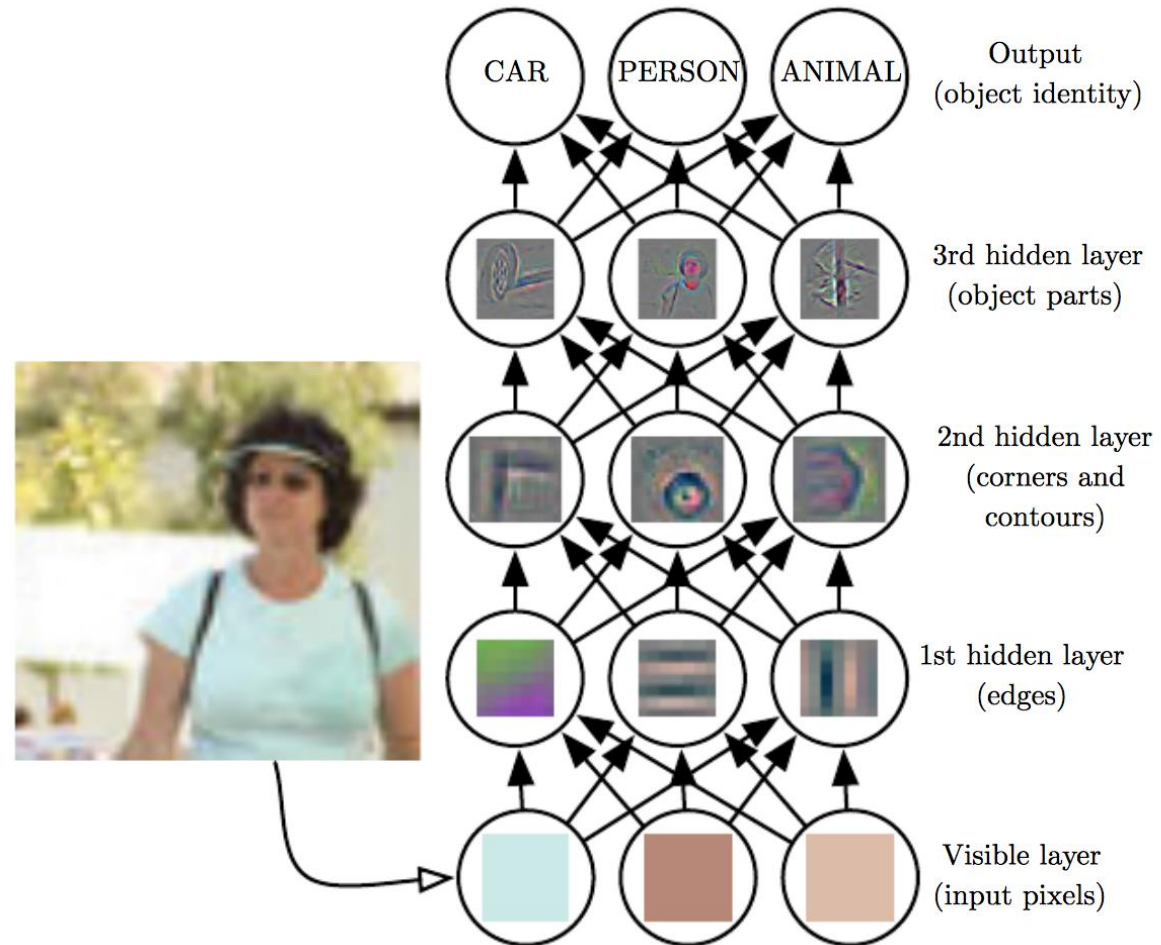
Pre-trained Imagenet models have good generalization ability.

Can be used for:

Prediction(use pretrained model directly)

Feature Extraction(only replace and retrain the classifier on top of the ConvNet )

Fine-tuning(fine-tune all the layers of the ConvNet, or keep some of the earlier layers fixed and fine-tune some higher-level layers)



# Table of contents

- Pre-training in image processing
- Language model to word embedding
- Word embedding to ELMO
- Word embedding to GTP
- The birth of BERT

# Language model

- The goal of language modelling is to estimate the probability distribution of various linguistic units, e.g., words, sentences etc
- Machine Translation:
  - $P(\text{high winds tonite}) > P(\text{large winds tonite})$
- Spell Correction
  - The office is about fifteen **minuets** from my house
    - $P(\text{about fifteen minutes from}) > P(\text{about fifteen minuets from})$
- Speech Recognition
  - $P(\text{I saw a van}) \gg P(\text{eyes awe of an})$
- Two categories: count-based and continuous-space LM.

# Probabilistic language model(count-based)

- Goal: compute the probability of a sentence or sequence of words:

$$P(W) = P(w_1, w_2, w_3, w_4, w_5 \dots w_n)$$

- Related task: probability of an upcoming word:


$$P(w_5 | w_1, w_2, w_3, w_4)$$

- A model that computes either of these:

$$P(W) \quad \text{or} \quad P(w_n | w_1, w_2 \dots w_{n-1}) \quad \text{is called a } \textbf{language model}.$$

Chain rule

$$P(S) = P(w_1, w_2, \dots, w_n) \xRightarrow{\text{Chain rule}} P(w_1, w_2, \dots, w_n) = P(w_1)P(w_2|w_1)P(w_3|w_1, w_2) \cdots P(w_n|w_1, w_2, \dots, w_{n-1})$$



$$L = \sum_{w \in C} \log P(w | \text{context}(w))$$



# Probabilistic language model(count-based)

Example: N-gram

- The LM probability  $p(w_1, w_2, \dots, w_n)$  is a product of word probabilities based on a history of  $m$  preceding words:

$$p(w_n | w_1, w_2, \dots, w_{n-1}) \approx p(w_n | w_{n-m}, \dots, w_{n-2}, w_{n-1})$$

- The estimation of a trigram word prediction probability:

$$p(w_3 | w_1, w_2) = \frac{\text{count}(w_1, w_2, w_3)}{\sum_w \text{count}(w_1, w_2, w)}$$

- Drawbacks:
  - Sparsity->smoothing
  - curse of dimensionality
  - rely on exact pattern(not linguistically informed)
  - dependency beyond the window is ignored

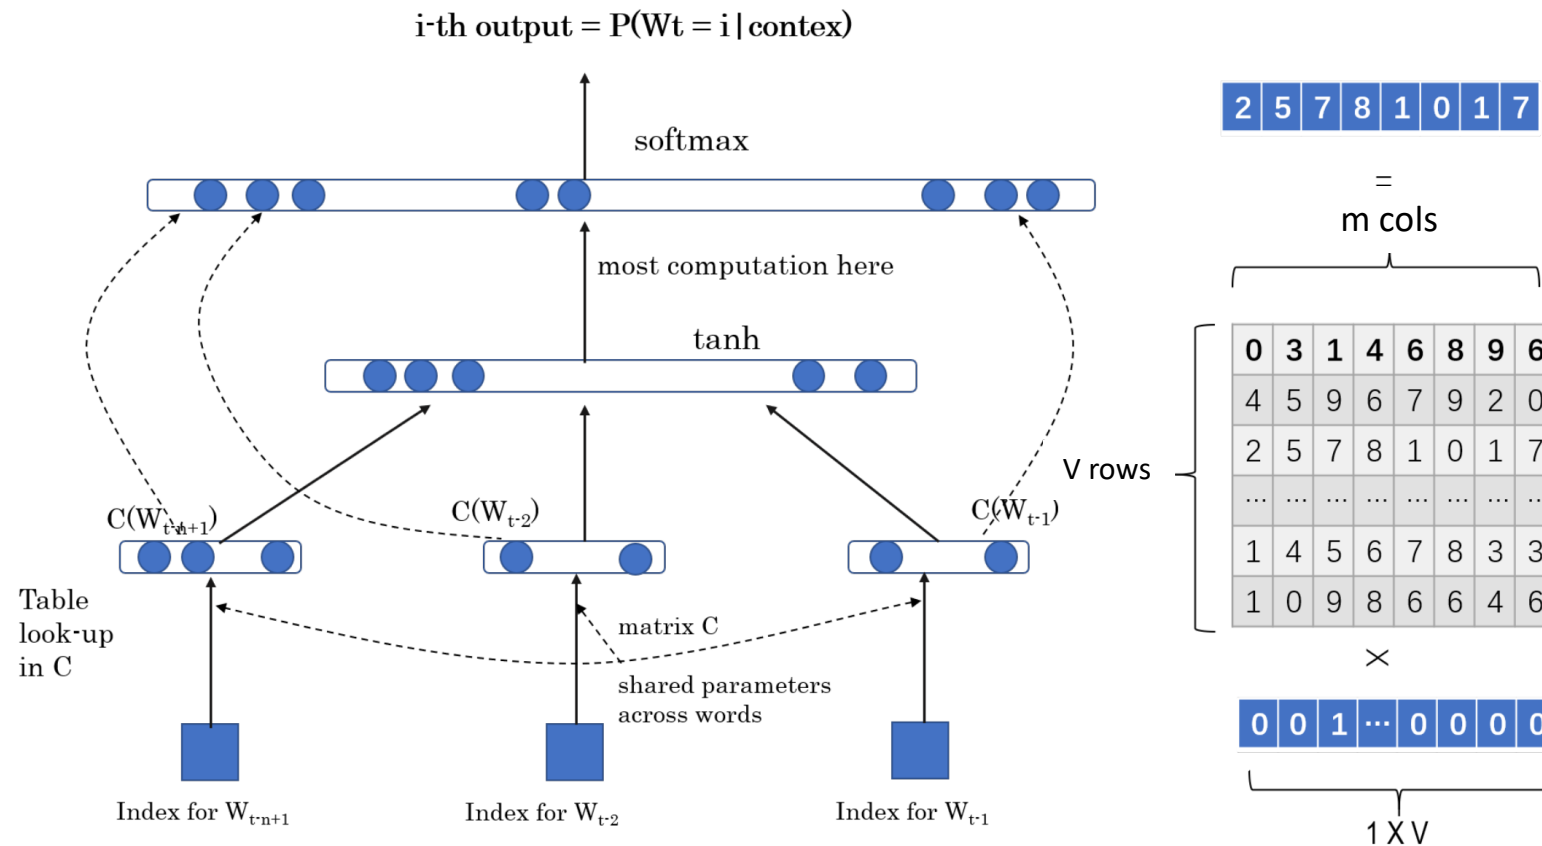
# Neural language models

## Two main forms:

- Feed-forward neural network based LM
  - to tackle the problems of data sparsity
- recurrent neural network based LM
  - to address the problem of limited context.


# Neural language model(Continuous-space)

- Feed-Forward Neural Network Based Models




# word2vec

A framework for learning word vectors

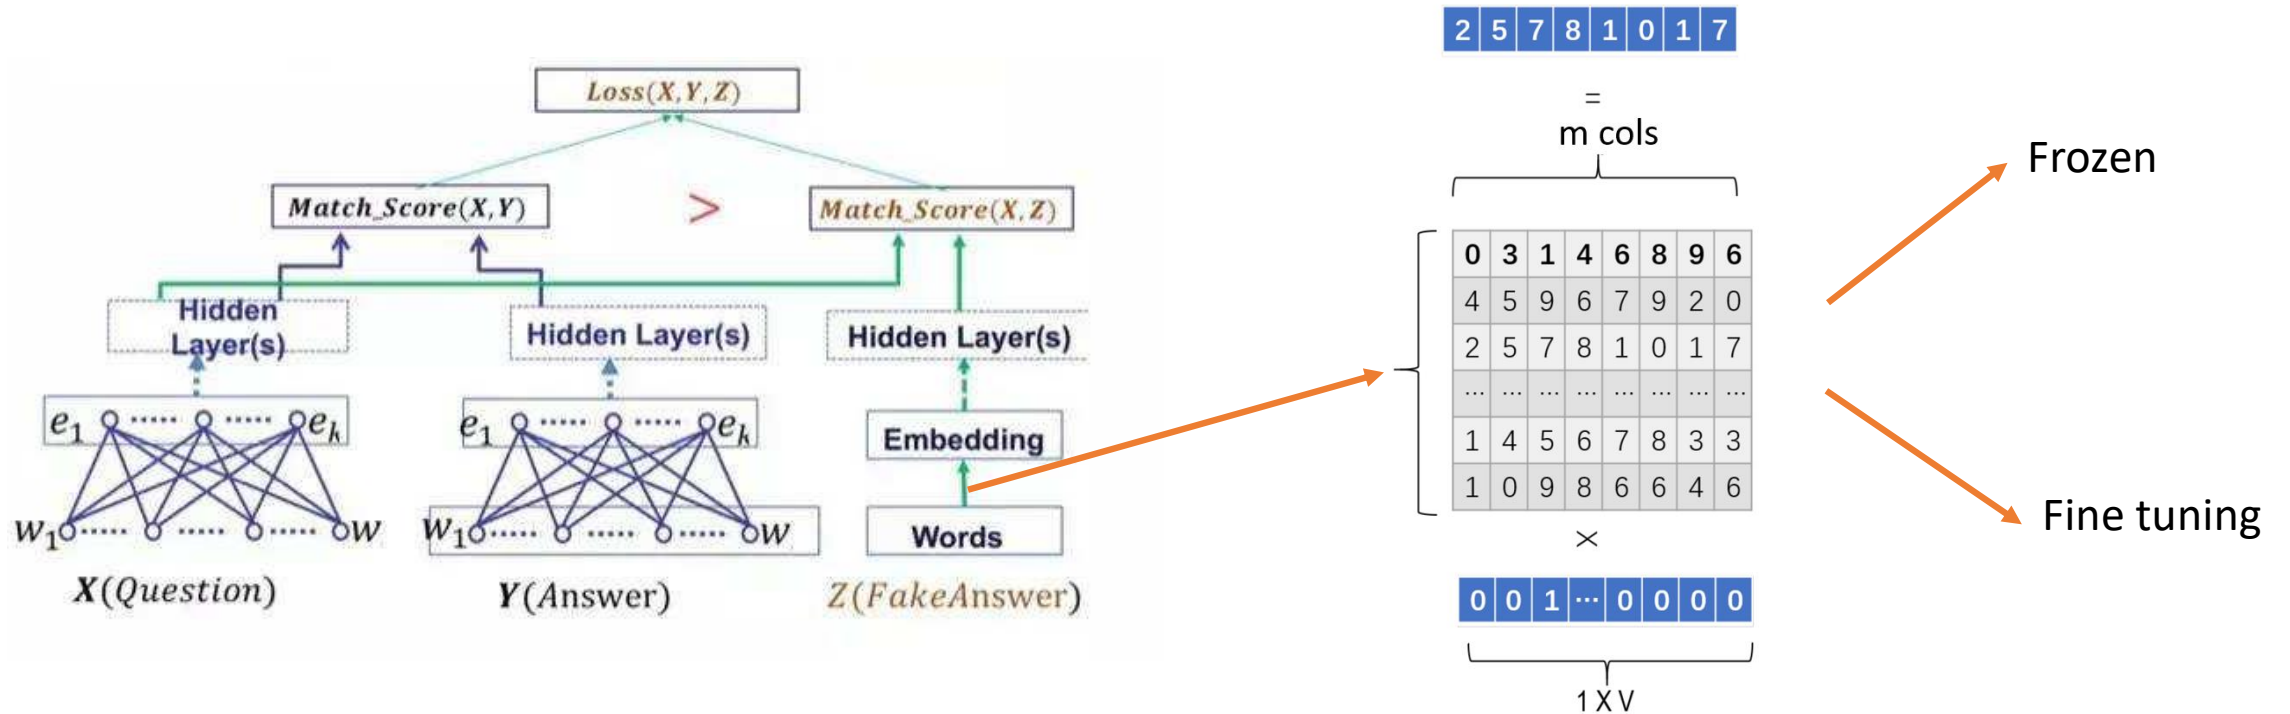


CBOW (Continuous Bag-of-Words Model)



Skip-gram (Continuous Skip-gram Model)

# How to use word embedding?



QA

Word Embedding Matrix


Classical usage of pretraining before 2018

# Limit of word2vec

...very useful to protect **banks** or slopes from being washed away by river or rain...  
...the location because it was high, about 100 feet above the **bank** of river...  
...The **bank** has plan to branch throughout the country...  
...They throttled the watchman and robbed the **bank**

## Polysemy: Bank

1. ORGANIZATION
2. RAISED GROUND



## Static word embedding:

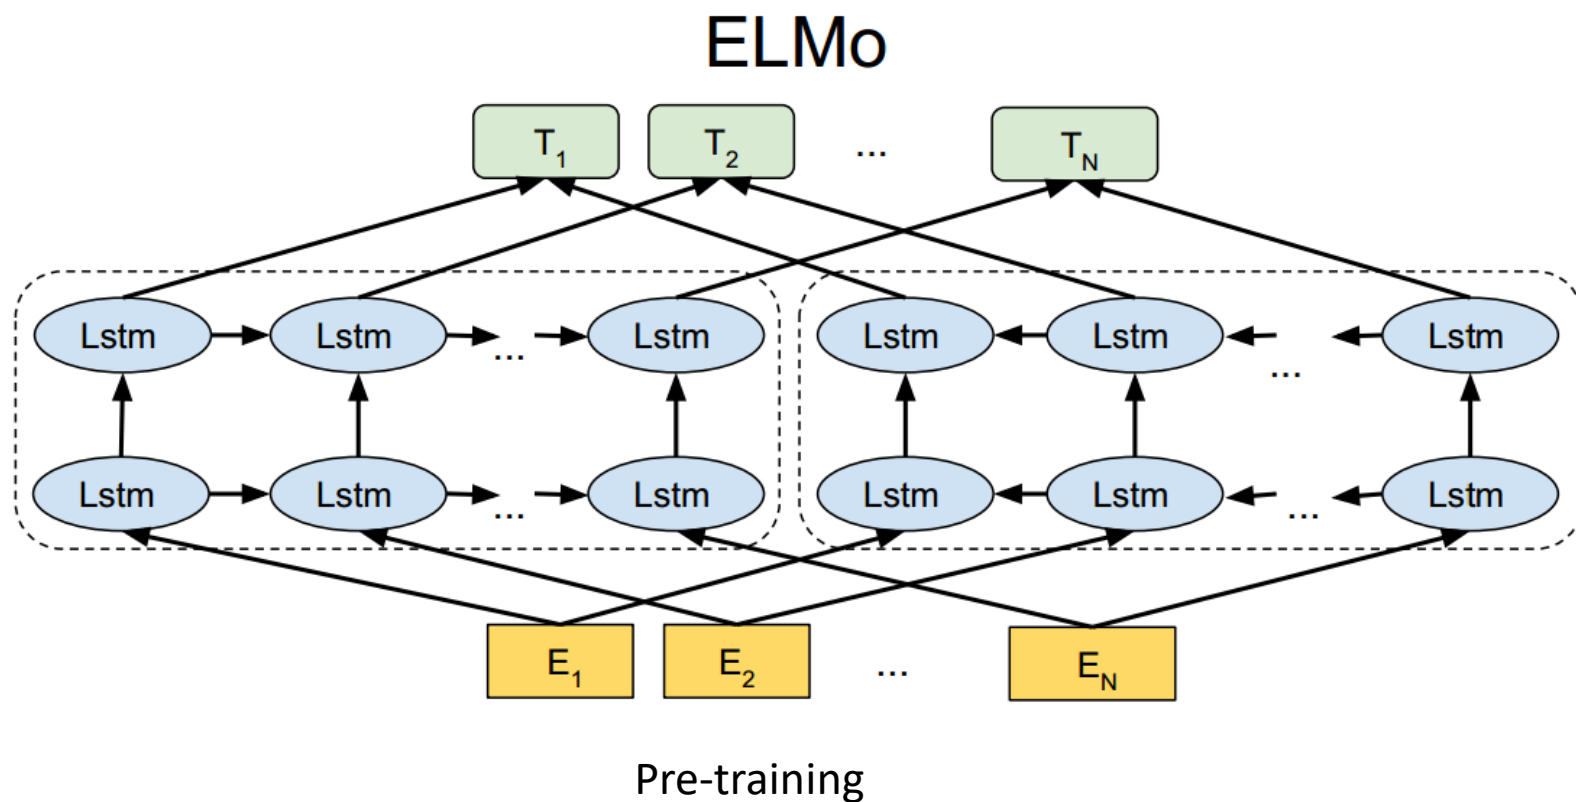
one vector for each word and  
smooshes all the context into that  
vector

# Table of contents

- Pre-training in image processing
- Language model to word embedding
- Word embedding to ELMO
- Word embedding to GTP
- The birth of BERT

# ELMo: Embeddings from Language Models (NAACL 2018)

Deep contextualized word representations

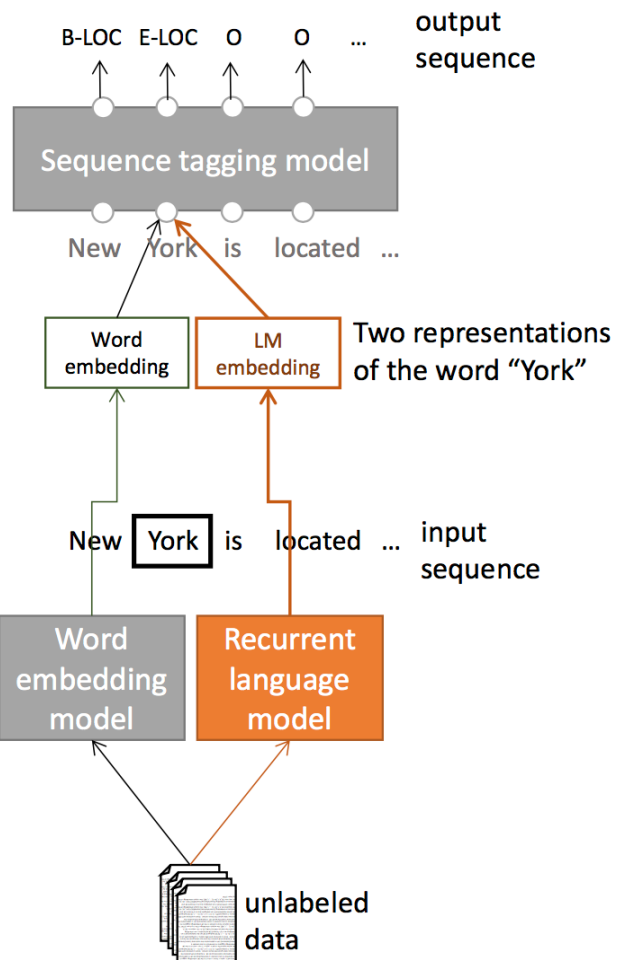




# ELMo: How to use?

## Step 3:

Use both word embeddings and LM embeddings in the sequence tagging model.



**Step 1:** Pretrain word embeddings and language model.

**Step 2:** Prepare word embedding and LM embedding for each token in the input sequence.

# ELMo: Polysemy?

Play:

1. exercise for recreation

2. drama

	Source	Nearest Neighbors
GloVe	play	playing, game, games, played, players, plays, player, Play, football, multiplayer
biLM	Chico Ruiz made a spectacular <u>play</u> on Alusik 's grounder {...}	Kieffer , the only junior in the group , was commended for his ability to hit in the clutch , as well as his all-round excellent <u>play</u> .
	Olivia De Havilland signed to do a Broadway <u>play</u> for Garson {...}	{...} they were actors who had been handed fat roles in a successful <u>play</u> , and had talent enough to fill the roles competently , with nice understatement .

Table 4: Nearest neighbors to “play” using GloVe and the context embeddings from a biLM.

- Solved!

# ELMo: Performance

TASK	PREVIOUS SOTA		OUR BASELINE	ELMo + BASELINE	INCREASE (ABSOLUTE/ RELATIVE)
SQuAD	<a href="#">Liu et al. (2017)</a>	84.4	81.1	85.8	4.7 / 24.9%
SNLI	<a href="#">Chen et al. (2017)</a>	88.6	88.0	$88.7 \pm 0.17$	0.7 / 5.8%
SRL	<a href="#">He et al. (2017)</a>	81.7	81.4	84.6	3.2 / 17.2%
Coref	<a href="#">Lee et al. (2017)</a>	67.2	67.2	70.4	3.2 / 9.8%
NER	<a href="#">Peters et al. (2017)</a>	$91.93 \pm 0.19$	90.15	$92.22 \pm 0.10$	2.06 / 21%
SST-5	<a href="#">McCann et al. (2017)</a>	53.7	51.4	$54.7 \pm 0.5$	3.3 / 6.8%

6 NLP tasks: improved by 5 ~ 25%

# ELMo: disadvantages

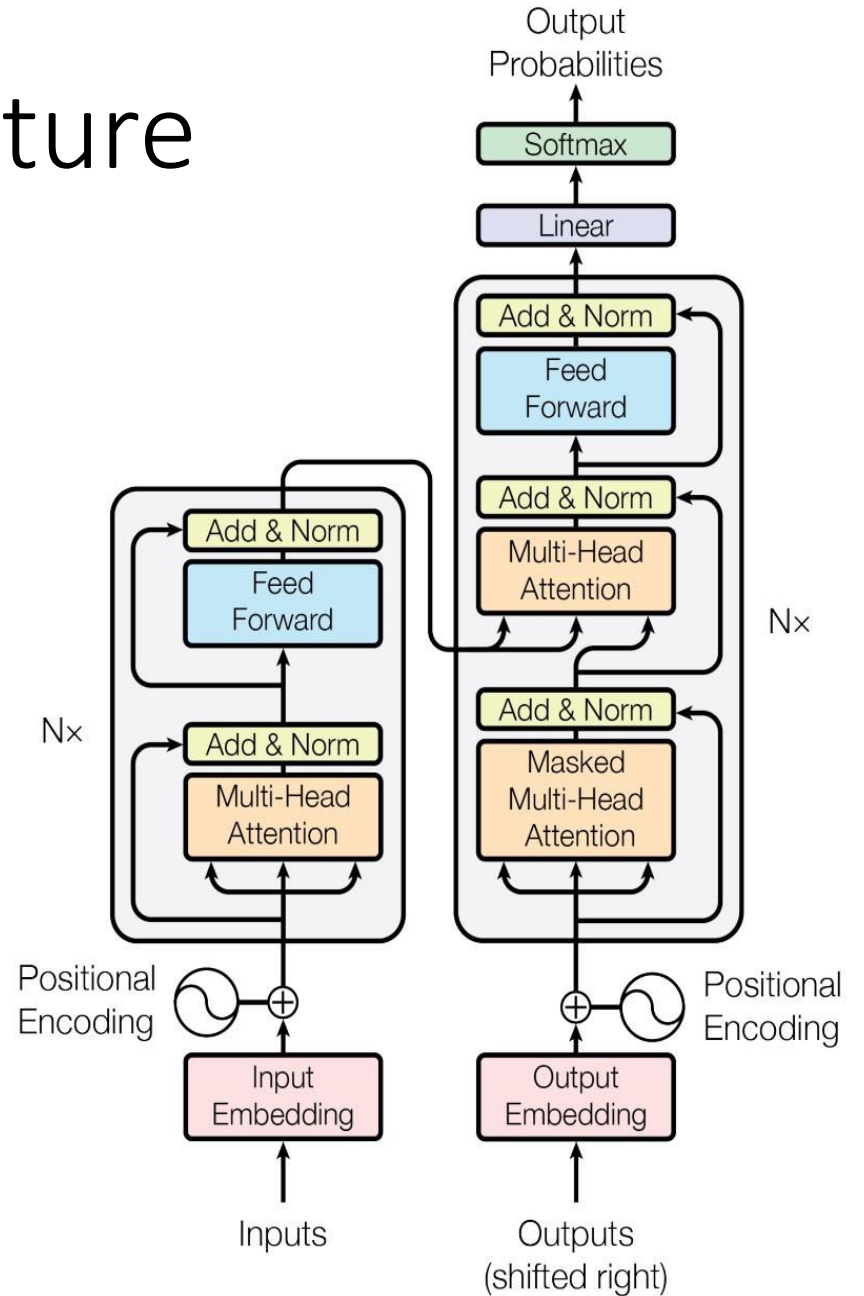
- LSTM is not as powerful as transformer in feature extraction
- Concatenation is not a ideal way to fuse the bi-directional information

# Table of contents

- Pre-training in image processing
- Language model to word embedding
- Word embedding to ELMO
- Word embedding to GTP
- The birth of BERT

# Transformer: overall architecture

- A sequence-to-sequence model based entirely on attention
- Strong results on standard WMT datasets
- Fast: only matrix multiplications




# Transformer: Attention Tricks

- **Self Attention:** Each layer combines words with Others
- **Multi-headed Attention:** 8 attention heads learned Independently
- **Normalized Dot-product Attention:** Remove bias in dot product when using large networks
- **Positional Encodings:** Make sure that even if we don't have RNN, can still distinguish positions

# Self Attention

## Scaled Dot-Product Attention



- Mapping a query and a set of key-value pairs to an output
- The output is a weighted sum of the values
- The weight assigned to each value is computed by a compatibility function of the query with the corresponding key



# Self Attention

Input

Embedding

Queries

Keys

Values

Score

Divide by 8 (  $\sqrt{d_k}$  )

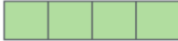
Softmax

Softmax

X  
Value

Sum

Thinking

$x_1$  

$q_1$  

$k_1$  

$v_1$  

$q_1 \cdot k_1 = 112$

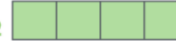
14

0.88

$v_1$  

$z_1$  

Machines

$x_2$  

$q_2$  

$k_2$  

$v_2$  

$q_1 \cdot k_2 = 96$


12

0.12

$v_2$  

$z_2$  

# Multi-Head Attention



- It expands the model's ability to focus on different positions.
- It gives the attention layer multiple representation subspaces


$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

# Multi-Head Attention



1) This is our input sentence\*

Thinking  
Machines


2) We embed each word\*




3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices



...





4) Calculate attention using the resulting  $Q/K/V$  matrices





...



5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^O$  to produce the output of the layer



...



\* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



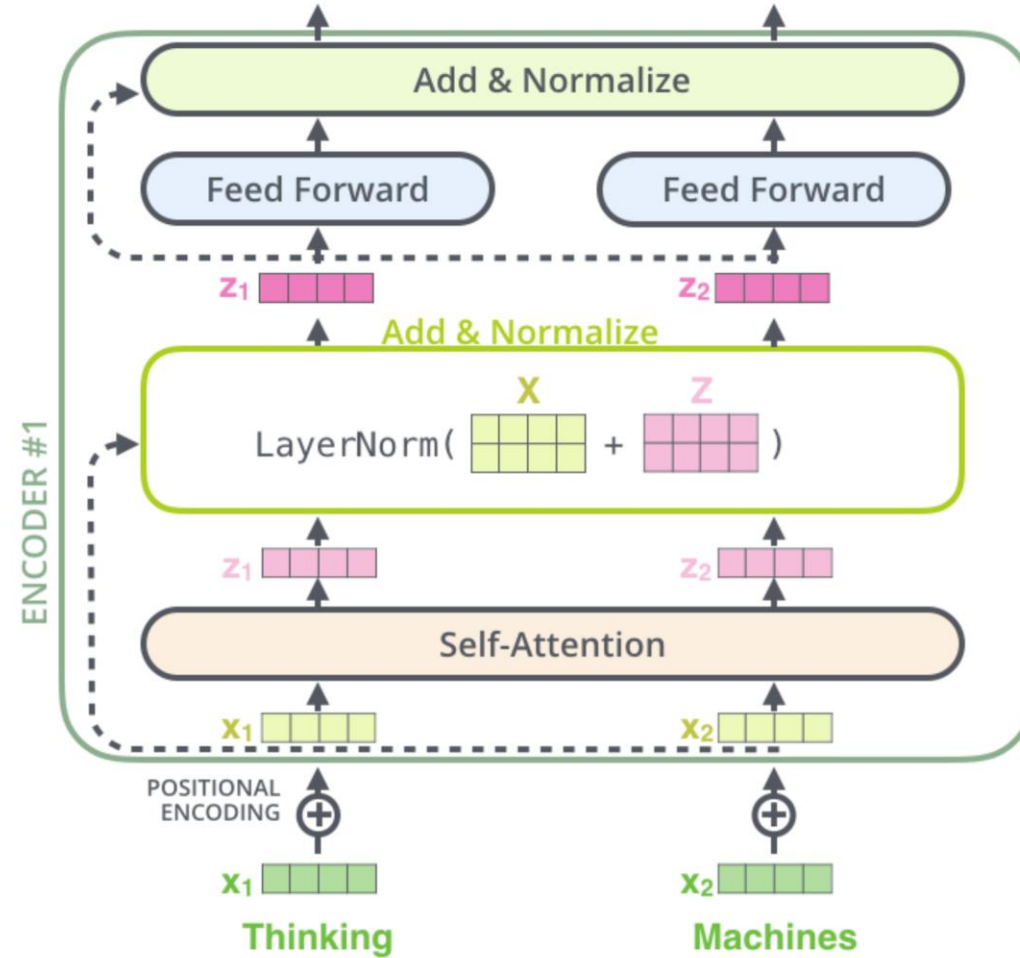
# Positional Encodings

- Encode the positions of the inputs as vectors and are then added to the input embeddings
- Each dimension of the positional encoding is a wave with a different frequency

$$PE[pos, 2i] = \sin(pos/10000^{2i/d_{model}})$$

$$PE[pos, 2i + 1] = \cos(pos/10000^{2i/d_{model}})$$

# Encoder

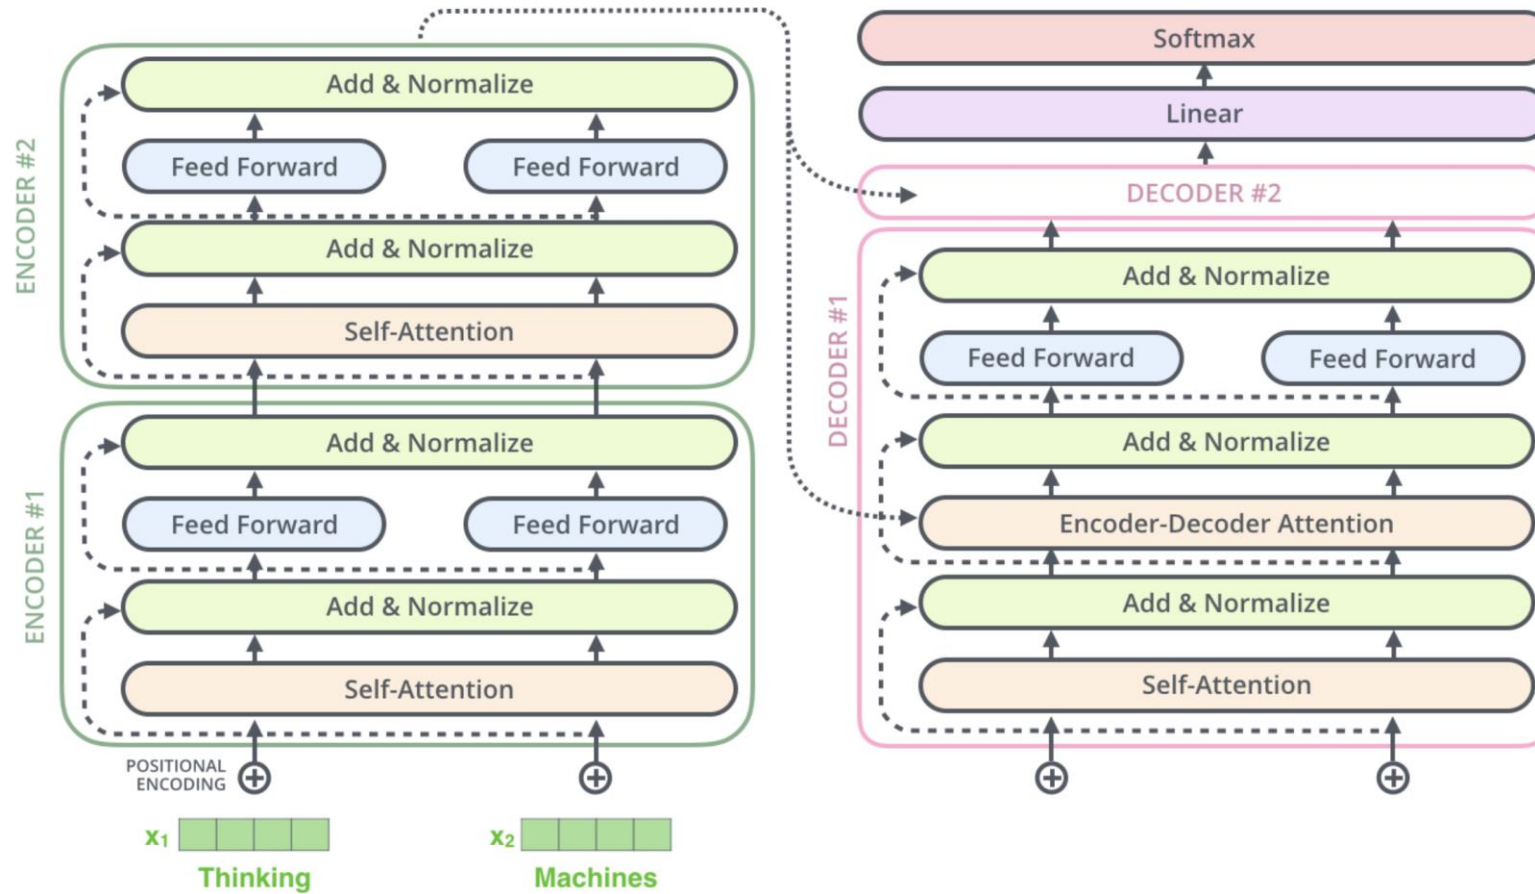


# Decoder

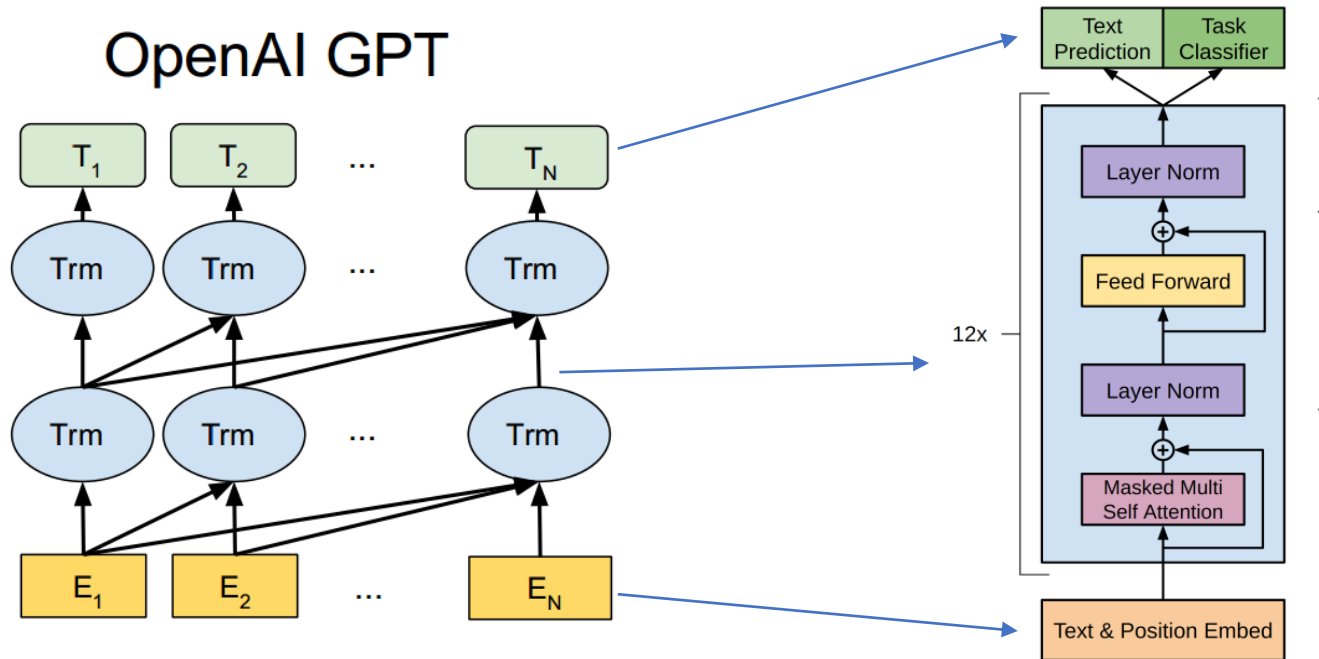
Masked multi-head self attention: self-attention layer is only allowed to attend to earlier positions in the output sequence

Encoder-Decoder Attention: attention layer works just like multiheaded self-attention, except it creates its Queries matrix from the layer previous to it, and takes the Keys and Values matrix from the output of the encoder stack

# Encoder-Decoder



# GPT: Generative Pre-Training (OpenAI 2018)

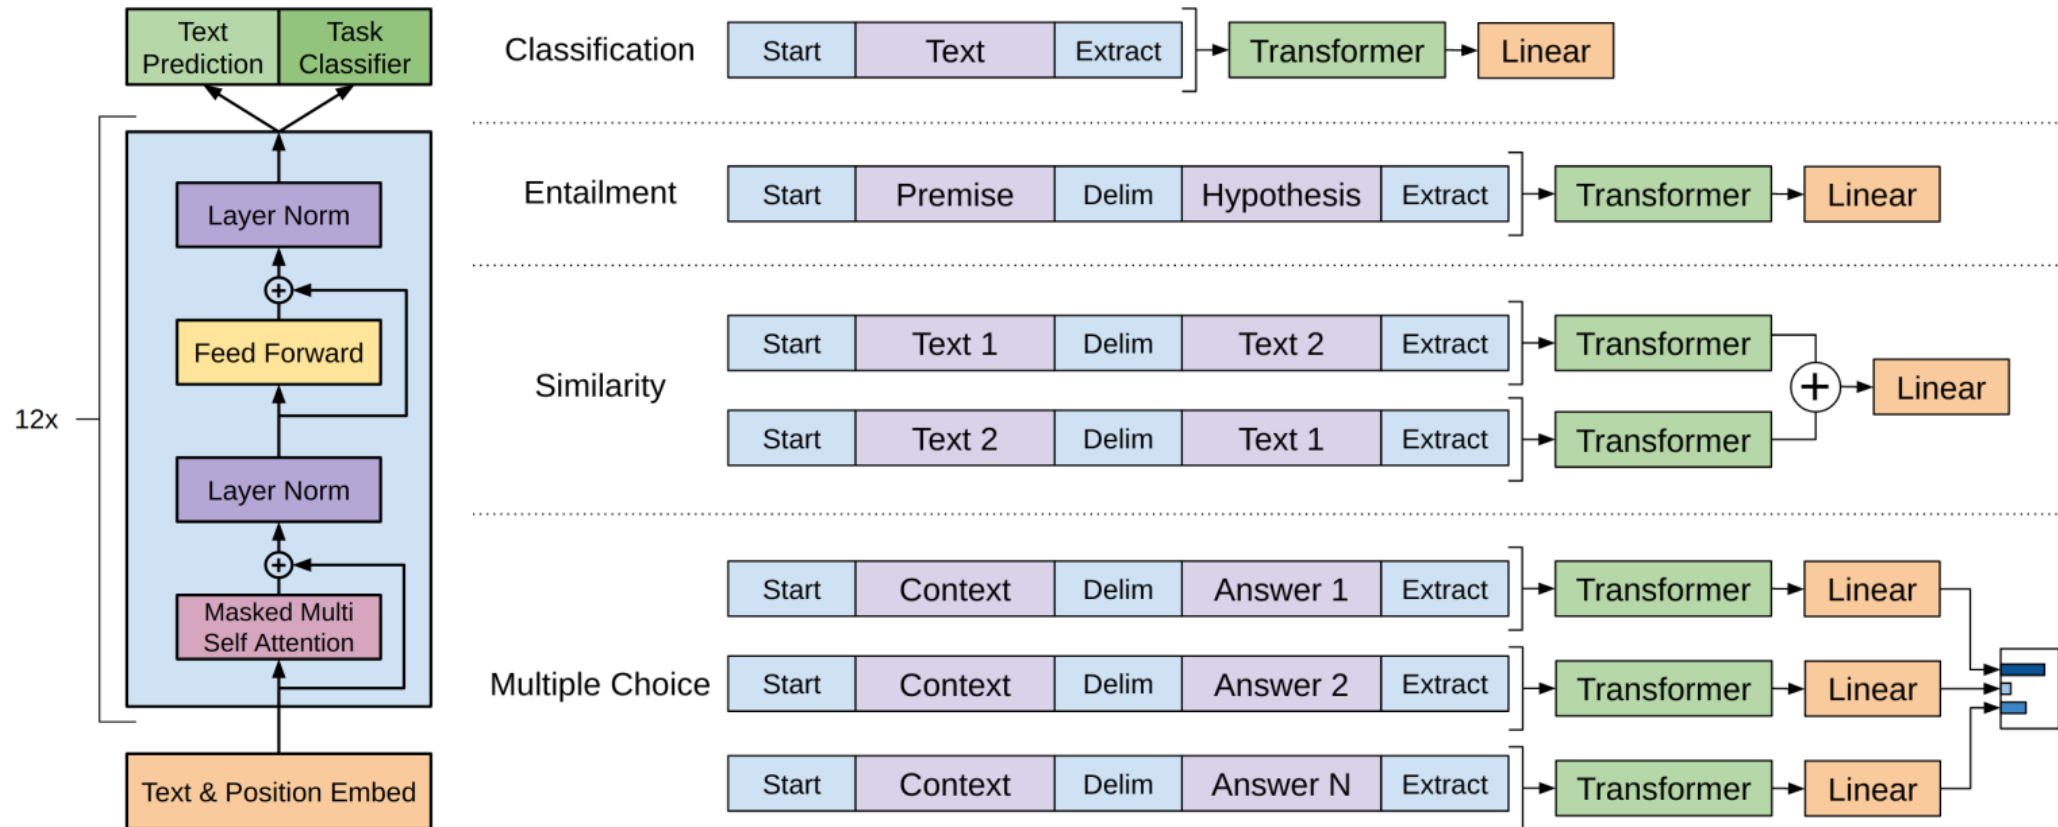


Pre-training

1. Transformer as feature extractor
2. pre-train a neural network using a language modeling objective (monodirectional)



# GPT: how to use?



# GPT: performance

Table 2: Experimental results on natural language inference tasks, comparing our model with current state-of-the-art methods. 5x indicates an ensemble of 5 models. All datasets use accuracy as the evaluation metric.

Method	MNLI-m	MNLI-mm	SNLI	SciTail	QNLI	RTE
ESIM + ELMo [44] (5x)	-	-	<u>89.3</u>	-	-	-
CAFE [58] (5x)	80.2	79.0	<u>89.3</u>	-	-	-
Stochastic Answer Network [35] (3x)	<u>80.6</u>	<u>80.1</u>	-	-	-	-
CAFE [58]	78.7	77.9	88.5	<u>83.3</u>		
GenSen [64]	71.4	71.3	-	-	<u>82.3</u>	59.2
Multi-task BiLSTM + Attn [64]	72.2	72.1	-	-	82.1	<b>61.7</b>
Finetuned Transformer LM (ours)	<b>82.1</b>	<b>81.4</b>	<b>89.9</b>	<b>88.3</b>	<b>88.1</b>	56.0

Achieve best performance in 9 of 12 NLP tasks

Table 3: Results on question answering and commonsense reasoning, comparing our model with current state-of-the-art methods.. 9x means an ensemble of 9 models.

Method	Story Cloze	RACE-m	RACE-h	RACE
val-LS-skip [55]	76.5	-	-	-
Hidden Coherence Model [7]	<u>77.6</u>	-	-	-
Dynamic Fusion Net [67] (9x)	-	55.6	49.4	51.2
BiAttention MRU [59] (9x)	-	<u>60.2</u>	<u>50.3</u>	<u>53.3</u>
Finetuned Transformer LM (ours)	<b>86.5</b>	<b>62.9</b>	<b>57.4</b>	<b>59.0</b>

Table 4: Semantic similarity and classification results, comparing our model with current state-of-the-art methods. All task evaluations in this table were done using the GLUE benchmark. (*mc*= Mathews correlation, *acc*=Accuracy, *pc*=Pearson correlation)

Method	Classification		Semantic Similarity			GLUE
	CoLA (mc)	SST2 (acc)	MRPC (F1)	STSB (pc)	QQP (F1)	
Sparse byte mLSTM [16]	-	<b>93.2</b>	-	-	-	-
TF-KLD [23]	-	-	<b>86.0</b>	-	-	-
ECNU (mixed ensemble) [60]	-	-	-	<u>81.0</u>	-	-
Single-task BiLSTM + ELMo + Attn [64]	<u>35.0</u>	90.2	80.2	55.5	<u>66.1</u>	64.8
Multi-task BiLSTM + ELMo + Attn [64]	18.9	91.6	83.5	72.8	<u>63.3</u>	<u>68.9</u>
Finetuned Transformer LM (ours)	<b>45.4</b>	91.3	82.3	<b>82.0</b>	<b>70.3</b>	<b>72.8</b>

# GPT: Analysis of various model ablations

Table 5: Analysis of various model ablations on different tasks. Avg. score is a unweighted average of all the results. (*mc*= Mathews correlation, *acc*=Accuracy, *pc*=Pearson correlation)

Method	Avg. Score	CoLA (mc)	SST2 (acc)	MRPC (F1)	STSB (pc)	QQP (F1)	MNLI (acc)	QNLI (acc)	RTE (acc)
Transformer w/ aux LM (full)	74.7	45.4	91.3	82.3	82.0	<b>70.3</b>	<b>81.8</b>	<b>88.1</b>	<b>56.0</b>
Transformer w/o pre-training	59.9	18.9	84.0	79.4	30.9	65.5	75.7	71.2	53.8
Transformer w/o aux LM	<b>75.0</b>	<b>47.9</b>	<b>92.0</b>	<b>84.9</b>	<b>83.2</b>	69.8	81.1	86.9	54.4
LSTM w/ aux LM	69.1	30.3	90.5	83.2	71.8	68.1	73.7	81.1	54.6

1. Transformer is better feature extractor than LSTM
2. Pre-training is critical

# GTP: disadvantage


- Monodirectional model

# Table of contents

- Pre-training in image processing
- Language model to word embedding
- Word embedding to ELMO
- Word embedding to GTP
- The birth of BERT

# BERT: Bidirectional Encoder Representations from Transformers(Google AI 2018)

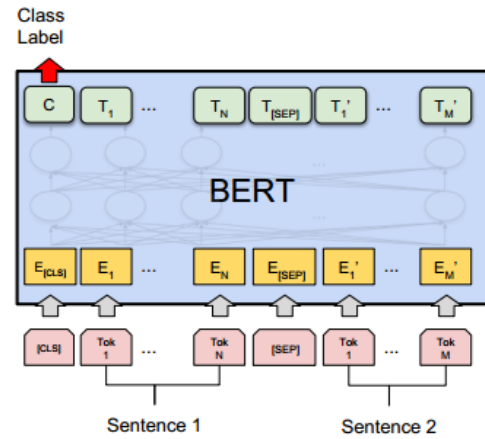
## BERT (Ours)



1. Transformer as feature extractor
2. Language modeling objective (bi-directional)

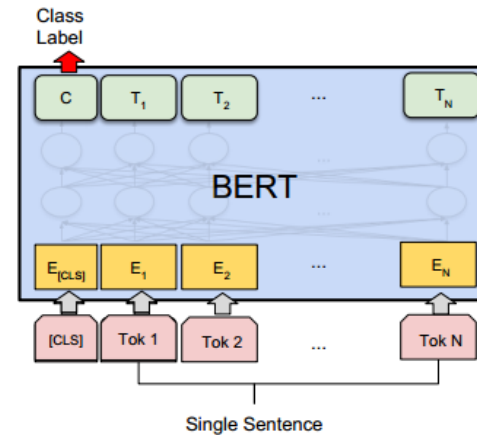
# BERT: how to use?

Relationship between sentences



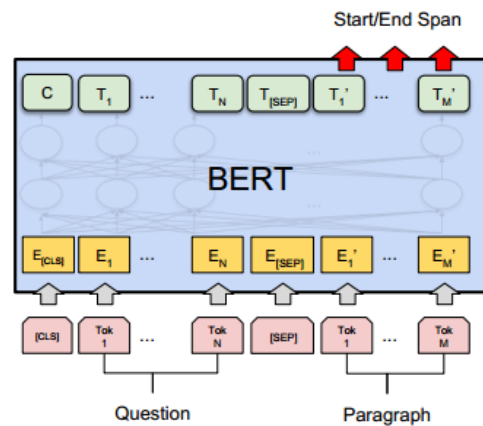
(a) Sentence Pair Classification Tasks:  
MNLI, QQP, QNLI, STS-B, MRPC,  
RTE, SWAG

Single sentence classification



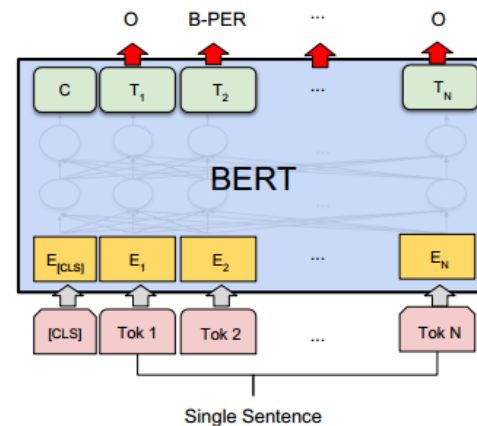
(b) Single Sentence Classification Tasks:  
SST-2, CoLA

Reading comprehension



(c) Question Answering Tasks:  
SQuAD v1.1

Sequence tagging



(d) Single Sentence Tagging Tasks:  
CoNLL-2003 NER

# BERT: performance

System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Average
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.9	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	88.1	91.3	45.4	80.0	82.3	56.0	75.2
BERT <sub>BASE</sub>	84.6/83.4	71.2	90.1	93.5	52.1	85.8	88.9	66.4	79.6
BERT <sub>LARGE</sub>	<b>86.7/85.9</b>	<b>72.1</b>	<b>91.1</b>	<b>94.9</b>	<b>60.5</b>	<b>86.5</b>	<b>89.3</b>	<b>70.1</b>	<b>81.9</b>

System	Dev F1	Test F1
ELMo+BiLSTM+CRF	95.7	92.2
CVT+Multi (Clark et al., 2018)	-	92.6
BERT <sub>BASE</sub>	96.4	92.4
BERT <sub>LARGE</sub>	<b>96.6</b>	<b>92.8</b>

Table 3: CoNLL-2003 Named Entity Recognition results. The hyperparameters were selected using the Dev set, and the reported Dev and Test scores are averaged over 5 random restarts using those hyperparameters.

System	Dev	Test
ESIM+GloVe	51.9	52.7
ESIM+ELMo	59.1	59.2
BERT <sub>BASE</sub>	81.6	-
BERT <sub>LARGE</sub>	<b>86.6</b>	<b>86.3</b>
Human (expert) <sup>†</sup>	-	85.0
Human (5 annotations) <sup>†</sup>	-	88.0

Table 4: SWAG Dev and Test accuracies. Test results were scored against the hidden labels by the SWAG authors. <sup>†</sup>Human performance is measure with 100 samples, as reported in the SWAG paper.

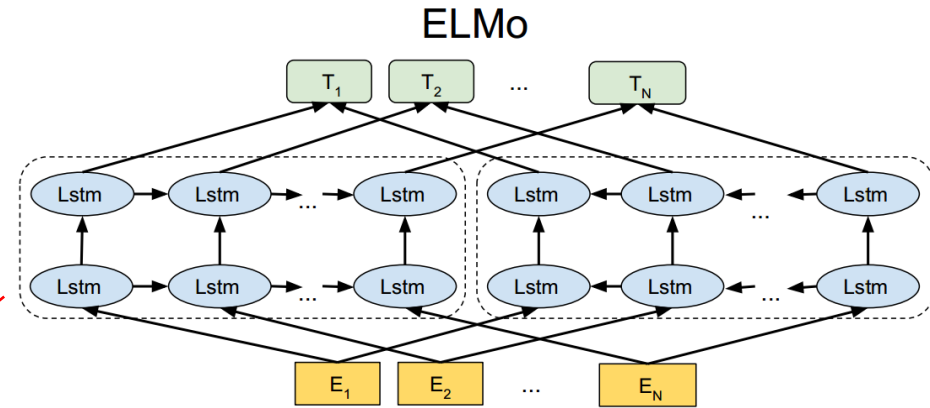

System	Dev		Test	
	EM	F1	EM	F1
Leaderboard (Oct 8th, 2018)				
Human	-	-	82.3	91.2
#1 Ensemble - nlnet	-	-	86.0	91.7
#2 Ensemble - QANet	-	-	84.5	90.5
#1 Single - nlnet	-	-	83.5	90.1
#2 Single - QANet	-	-	82.5	89.3
Published				
BiDAF+ELMo (Single)	-	85.8	-	-
R.M. Reader (Single)	78.9	86.3	79.5	86.6
R.M. Reader (Ensemble)	81.2	87.9	82.3	88.5
Ours				
BERT <sub>BASE</sub> (Single)	80.8	88.5	-	-
BERT <sub>LARGE</sub> (Single)	84.1	90.9	-	-
BERT <sub>LARGE</sub> (Ensemble)	85.8	91.8	-	-
BERT <sub>LARGE</sub> (Sgl.+TriviaQA)	<b>84.2</b>	<b>91.1</b>	<b>85.1</b>	<b>91.8</b>
BERT <sub>LARGE</sub> (Ens.+TriviaQA)	<b>86.2</b>	<b>92.2</b>	<b>87.4</b>	<b>93.2</b>

Table 2: SQuAD results. The BERT ensemble is 7x systems which use different pre-training checkpoints and fine-tuning seeds.

Improve the performance of all 11 tasks




# From word embedding to BERT




bidirectional LM

Transformer


BERT (Ours)




?




# BERT: how to construct bi-directional LM?



Similar idea



But want Trm architecture



Next sentence prediction

Masked LM

# BERT: how to construct bi-directional LM?

token. Instead, the training data generator chooses 15% of tokens at random, e.g., in the sentence `my dog is hairy` it chooses `hairy`. It then performs the following procedure:

- Rather than *always* replacing the chosen words with [MASK], the data generator will do the following:
- 80% of the time: Replace the word with the [MASK] token, e.g., `my dog is hairy` → `my dog is [MASK]`
- 10% of the time: Replace the word with a random word, e.g., `my dog is hairy` → `my dog is apple`
- 10% of the time: Keep the word unchanged, e.g., `my dog is hairy` → `my dog is hairy`. The purpose of this is to bias the representation towards the actual observed word.



Masked LM

# BERT: how to construct bi-directional LM?

not directly captured by language modeling. In order to train a model that understands sentence relationships, we pre-train a binarized *next sentence prediction* task that can be trivially generated from any monolingual corpus. Specifically, when choosing the sentences A and B for each pre-training example, 50% of the time B is the actual next sentence that follows A, and 50% of the time it is a random sentence from the corpus. For example:

**Input** = [CLS] the man went to [MASK] store [SEP]

he bought a gallon [MASK] milk [SEP]

**Label** = IsNext

**Input** = [CLS] the man [MASK] to the store [SEP]

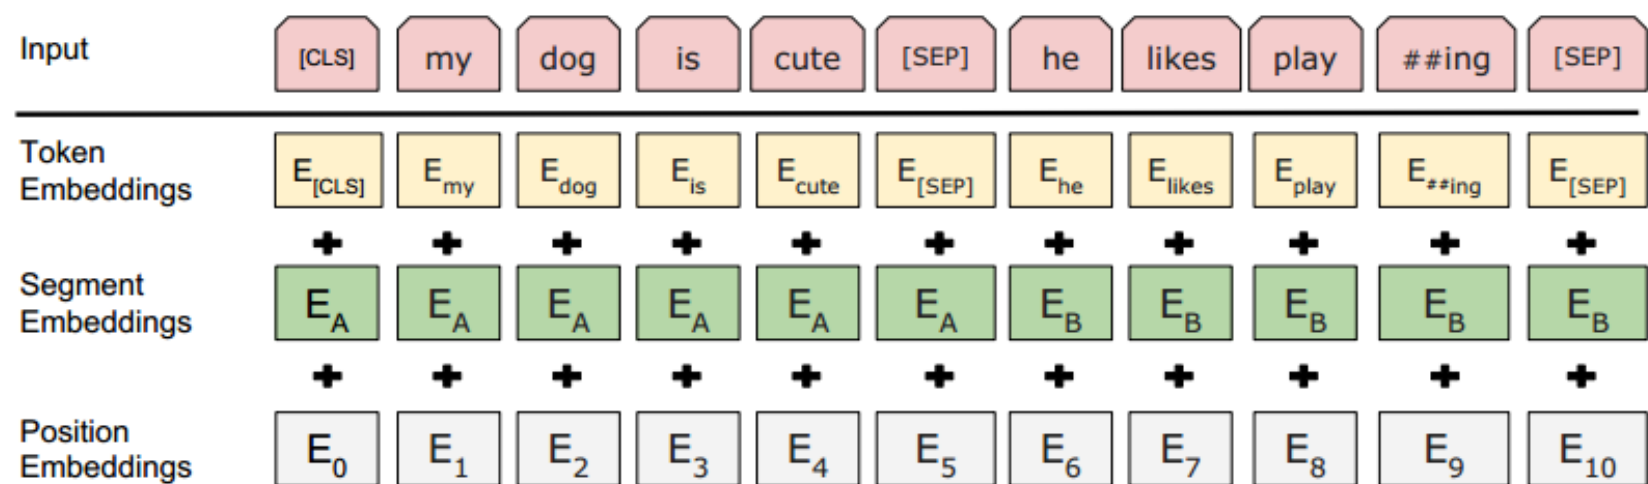
penguin [MASK] are flight ##less birds [SEP]

**Label** = NotNext



Next sentence prediction

# BERT: input processing



# BERT: Ablation analysis

Tasks	Dev Set				
	MNLI-m (Acc)	QNLI (Acc)	MRPC (Acc)	SST-2 (Acc)	SQuAD (F1)
BERT <sub>BASE</sub>	84.4	88.4	86.7	92.7	88.5
No NSP	83.9	84.9	86.5	92.6	87.9
LTR & No NSP	82.1	84.3	77.5	92.1	77.8
+ BiLSTM	82.1	84.1	75.7	91.6	84.9

1. Bidirectional language model is important
2. Next sentence prediction is critical for certain tasks

Table 5: Ablation over the pre-training tasks using the BERT<sub>BASE</sub> architecture. “No NSP” is trained without the next sentence prediction task. “LTR & No NSP” is trained as a left-to-right LM without the next sentence prediction, like OpenAI GPT. “+ BiLSTM” adds a randomly initialized BiLSTM on top of the “LTR + No NSP” model during fine-tuning.