

---

# **AlexNet, VGG, GoogleNet and ResNet**

---

**Presented by Zhengxia Zou**  
**18 Aug. 2018**

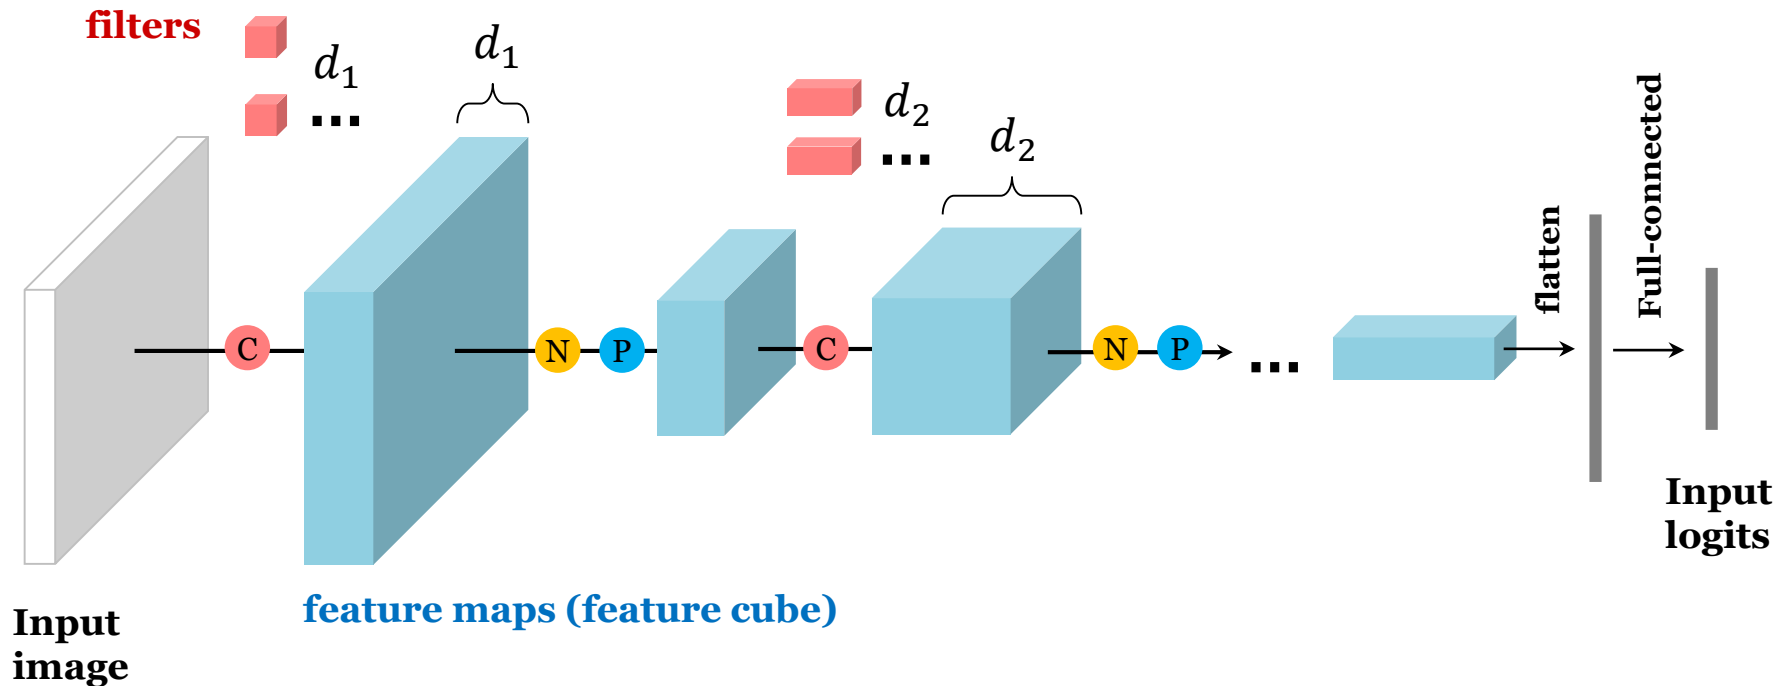
- 
- **Basic components of CNN**
  - **AlexNet**
  - **VGG**
  - **GoogleNet**
  - **ResNet**

# Basic components of CNN

**C** : Convolutional layer

**N** : Nonlinear mapping layer

**P** : Pooling layer





# Basic components of CNN

**C** : Convolutional layer



**N** : Nonlinear mapping layer

**P** : Pooling layer



# LeNet-5

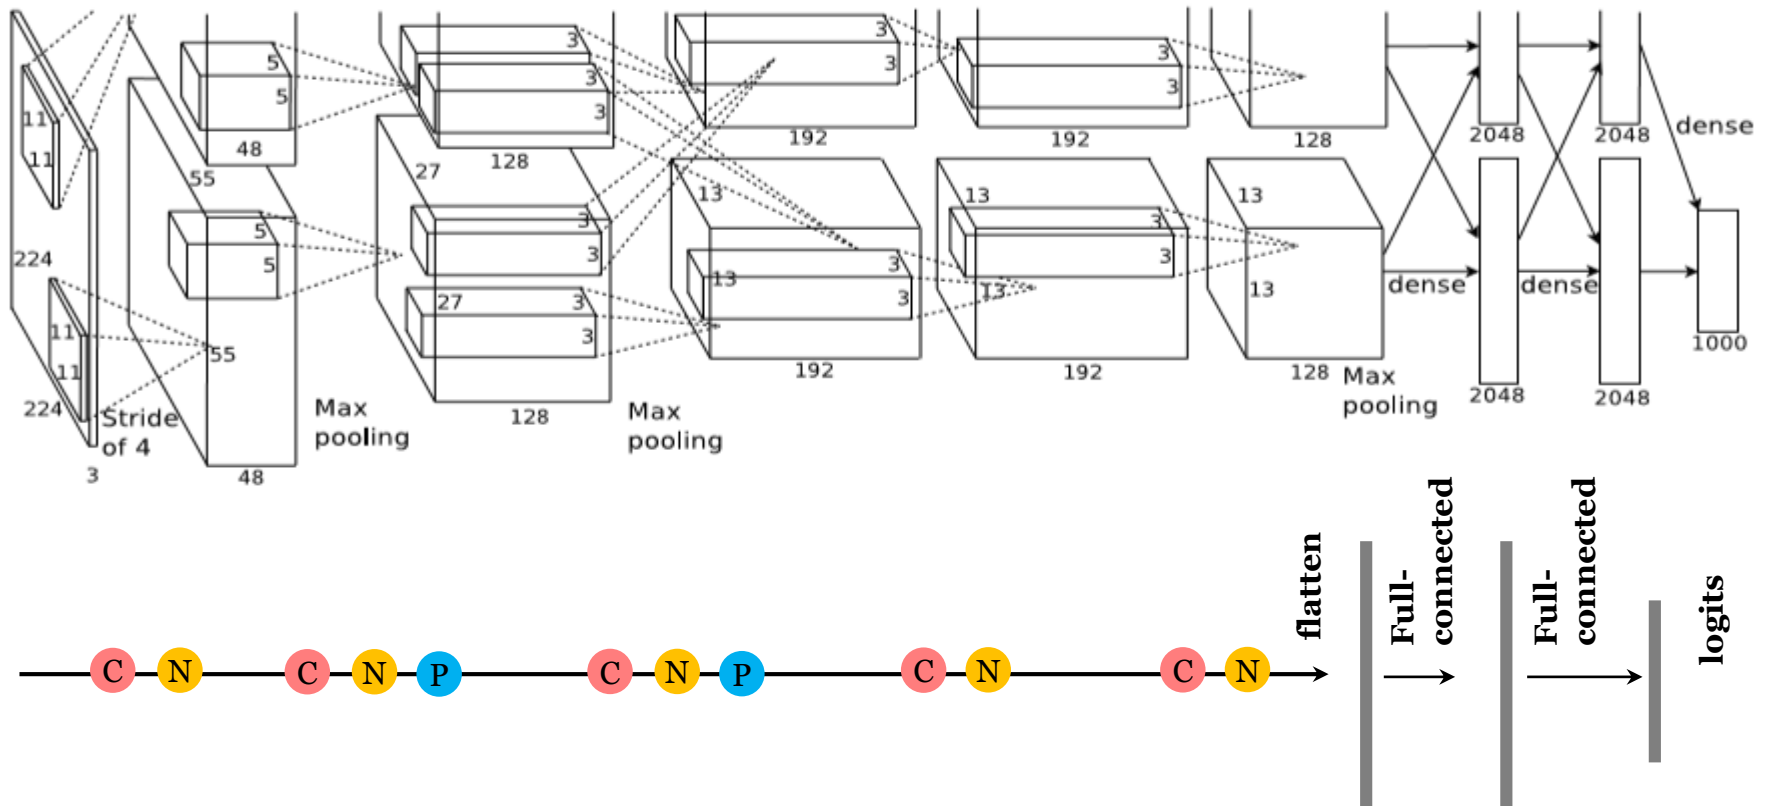
Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.



# AlexNet

Krizhevsky, Alex, I. Sutskever, and G. E. Hinton. "ImageNet classification with deep convolutional neural networks." *NIPS*, 2012.

2012 ImageNet top-5 test error: 1<sup>st</sup>: 15.3%, 2<sup>nd</sup>: 26.2%




# AlexNet


---

## Key points of AlexNet

- Rectified Linear Unit (ReLU)
- Dropout
- Local Response Normalization (LRN)



ReLU





Dropout

# VGG

Karen Simonyan, Andrew Zisserman. “Very Deep Convolutional Networks for Large-Scale Image Recognition.” *ICLR*, 2015. (arXiv 2014)

@<https://www.cs.toronto.edu/~frossard/post/vgg16/>






# VGG

---

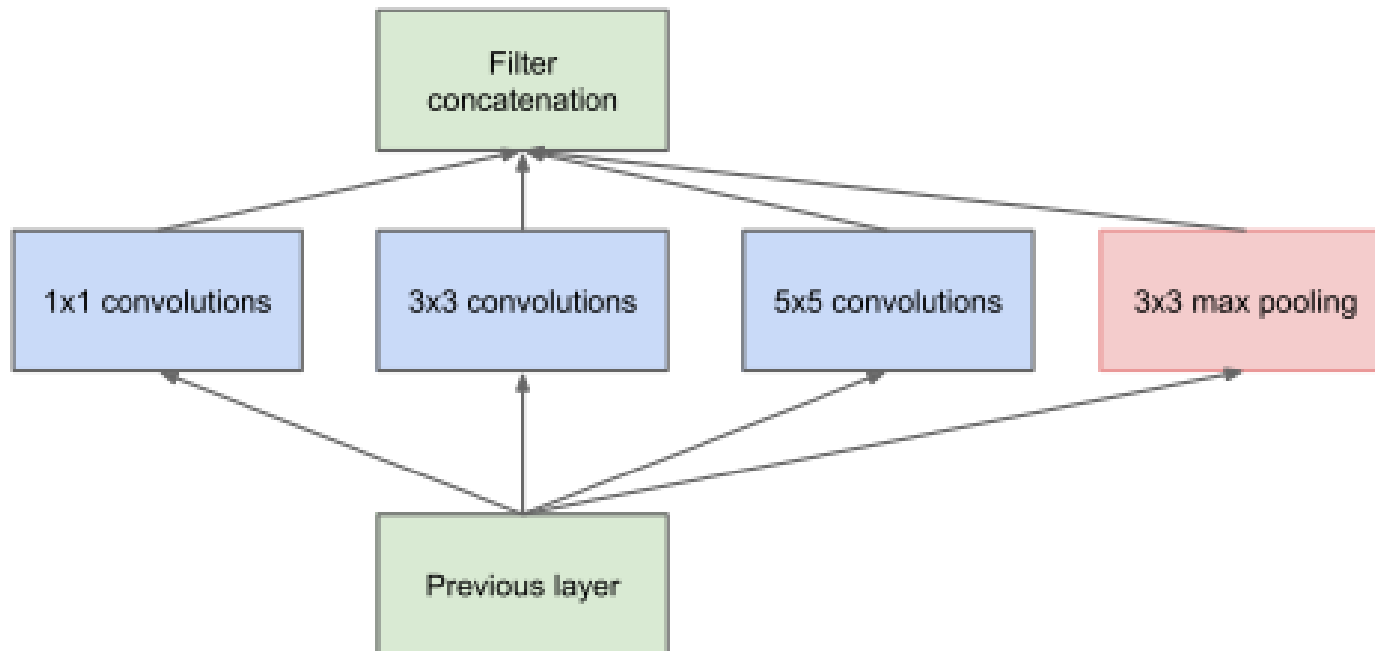
## Key points of VGG

- Increasing depth (16-19 layers)
- Using very small (3x3) convolution filters (instead of 5x5 and 7x7)



# GoogleNet

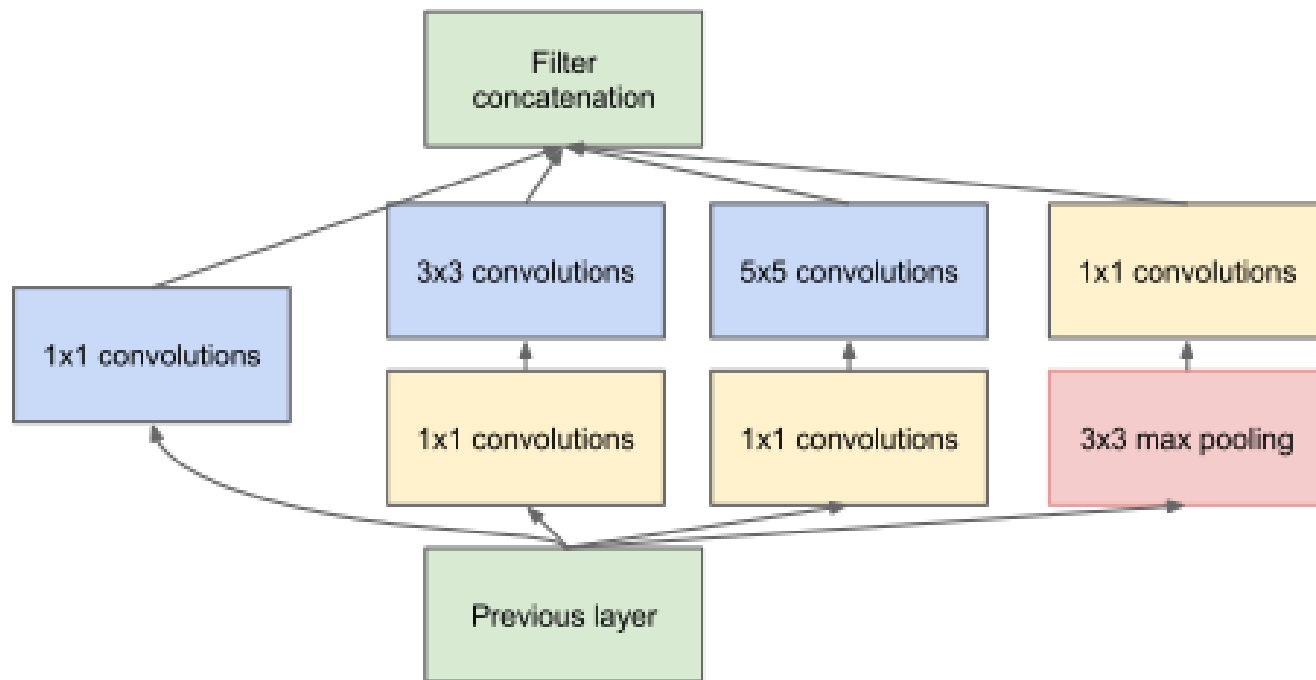
Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, et, al. “Going Deeper with Convolutions.” *CVPR*, 2015. (arXiv 2014)



(a) Inception module, naïve version

# GoogleNet

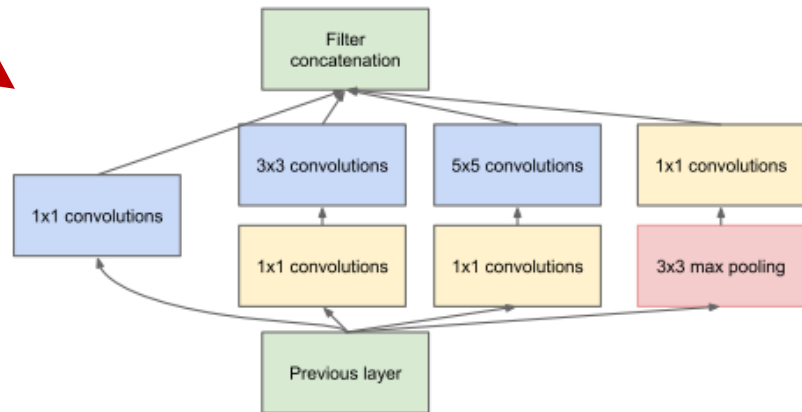
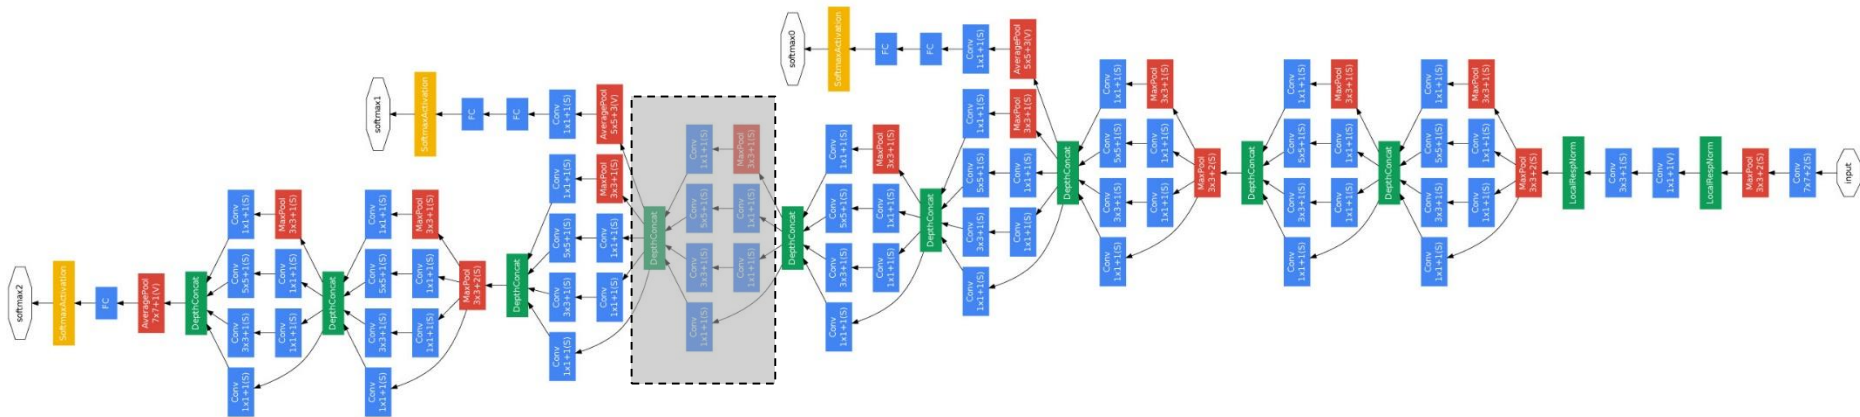
Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, et, al. “Going Deeper with Convolutions.” *CVPR*, 2015. (arXiv 2014)



(b) Inception module with dimensionality reduction

# GoogleNet

Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, et, al. “Going Deeper with Convolutions.” *CVPR*, 2015. (arXiv 2014)



(b) Inception module with dimensionality reduction

# GoogleNet

---

## Key points of GoogleNet (Inception v1-v4)

- Multiscale convolutions
- Increasing both of the width and depth (22 layers)
- Batch Normalization (BN)

$$\hat{x}^{(k)} = \frac{x^k - E[x^k]}{\sqrt{Var[x^k]}}$$

$$y^{(k)} = \gamma^k \hat{x}^{(k)} + \beta^{(k)}$$



Act as a pre-processing step:  
shifting inputs of each layer to zero-mean and  
unit variance



# ResNet

Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun. “Deep Residual Learning for Image Recognition.” *CVPR*, 2016, best paper. (arXiv 2015)

## Motivation


To ease the training of networks that are much deeper than those used previously.

Results on CIFAR10




# ResNet

## Building Block of ResNet



naïve version



“bottleneck” version

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x}.$$


# ResNet

---

**C** : Convolutional layer

**N** : Nonlinear mapping layer

**P** : Pooling layer





# ResNet

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	$112 \times 112$	$7 \times 7, 64, \text{stride } 2$				
conv2_x	$56 \times 56$	$3 \times 3 \text{ max pool, stride } 2$				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	$28 \times 28$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	$14 \times 14$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	$7 \times 7$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	$1 \times 1$	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# ResNet

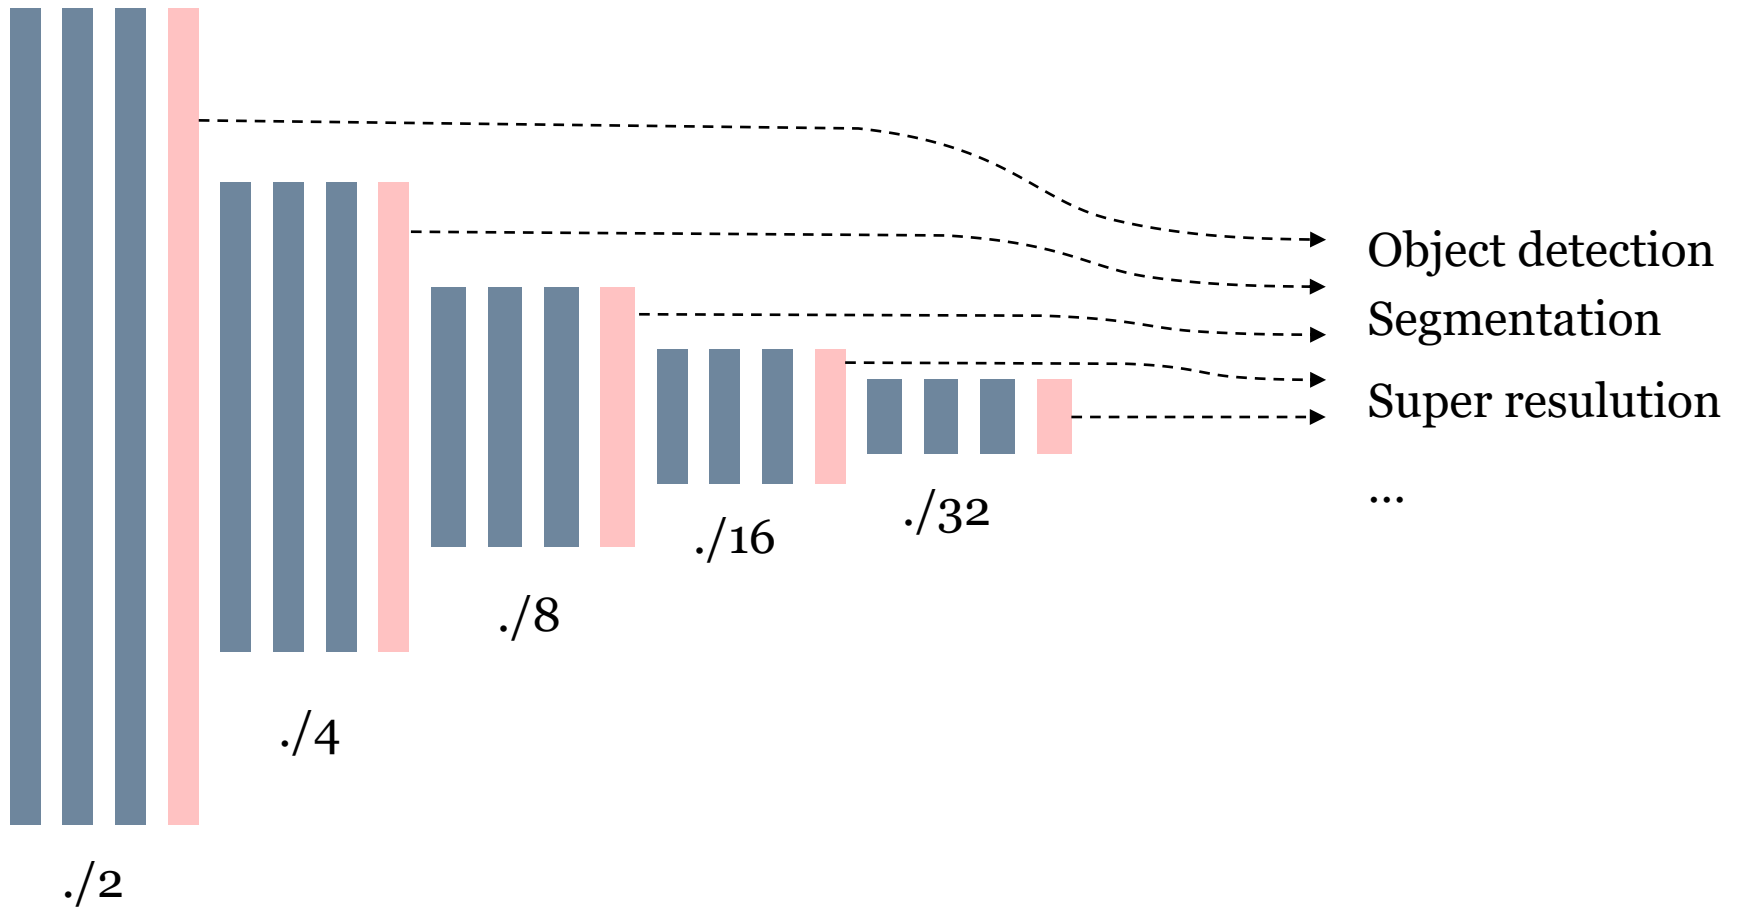
---

## Key points of ResNet

- Learning Residuals
- Much easier to optimize
- Gain accuracy from considerably increased depth (152 layers)

# Apply to other tasks

---





# Comparison

@ Alfredo Canziani, Adam Paszke, Eugenio Culurciello. An Analysis of Deep Neural Network Models for Practical Applications. arXiv 2016.

