

Deep Reinforcement Learning

Wenjun Zeng

Reading Group of Ye Lab, UM

Mar. 9, 2019

Contents

- Recap
- Q -Learning
- Deep Q -Networks (DQN)

I. Recap

Markov Decision Processes (MDP)

- **State space** \mathcal{S} , $s \in \mathcal{S}$
- **Action space** \mathcal{A} , $a \in \mathcal{A}$
- **Policy** π , $a = \pi(s)$ for deterministic case or $a \sim \pi(a|s)$ for stochastic case
- **Reward** $R(s, a) : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$
- **Transition Dynamics** $P: \mathcal{S} \times \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}_+$
 $P_{sa}(s')$ is the probability that the agent reaches state s' from s after taking action a .

Objective of MDP

Denoting s_t , a_t , and $R_t = R(s_t, a_t)$ as the state, action, and reward of $t = 0, 1, 2, \dots$, MDP aims at finding policy to maximize the expectation of

$$G_t = \sum_{t=0}^{\infty} \gamma^t R_t$$

with $\gamma \in [0, 1)$ the **discount factor**.

□ Model-free: State transition $P_{sa}(s')$ and reward function are unknown.

Q-function (state-action value function)

$$Q^{\pi}(s, a) = \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) \middle| s_0 = s, a_0 = a, \pi \right]$$

is the expected return for executing a particular action at a given state.

Bellman's Equation

- Bellman equation on Q -function

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{sa}(s') \max_{a' \in \mathcal{A}} Q^*(s', a')$$

- Bellman equation in expectation form:

$$Q^*(s, a) = \mathbb{E}_{s' \sim P_{sa}(s')} \left[R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q^*(s', a') \right]$$

- In conventional optimization (linear/nonlinear programming...), we solve **Karush–Kuhn–Tucker (KKT) condition** to obtain the optimal (stationary) solution.
- In MDP (RL), we solve **Bellman's equation** to obtain the optimal solution. If $P_{sa}(s')$ is known (model-based), it is dynamic programming.

Q -Value Iteration for Model-Based Case

□ $Q^*(s, a)$: expected return starting in s , taking action a , and (thereafter) acting optimally

□ Bellman Equation:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} P_{sa}(s') (R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q^*(s', a'))$$

□ Q -Value Iteration:

$$Q(s, a) \leftarrow \sum_{s' \in \mathcal{S}} P_{sa}(s') (R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q(s', a'))$$

Model-Free Reinforcement Learning

- In realistic problems, often the state transition and reward function are not explicitly given.
- Model-free RL is to **directly learn** value & policy from **experience**.
- How to accumulate experience? → learning from **episodes**

For $t = 0, 1, 2, \dots, T$

$$\text{Episode 1: } s_0^{(1)} \xrightarrow[R(s_0)^{(1)}]{a_0^{(1)}} s_1^{(1)} \xrightarrow[R(s_1)^{(1)}]{a_1^{(1)}} s_2^{(1)} \xrightarrow[R(s_2)^{(1)}]{a_2^{(1)}} s_3^{(1)} \dots s_T^{(1)}$$

$$\text{Episode 2: } s_0^{(2)} \xrightarrow[R(s_0)^{(2)}]{a_0^{(2)}} s_1^{(2)} \xrightarrow[R(s_1)^{(2)}]{a_1^{(2)}} s_2^{(2)} \xrightarrow[R(s_2)^{(2)}]{a_2^{(2)}} s_3^{(2)} \dots s_T^{(2)}$$

(Tabular) Q -Learning

Rewrite Q -value iteration as expectation:

$$Q(s, a) \leftarrow \mathbb{E}_{s' \sim P_{sa}(s')} \left[R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q(s', a') \right]$$

(Tabular) Q -Learning: replace expectation by samples

□ For a state-action pair (s, a) , receive: $s' \sim P_{sa}(s')$

□ Use the old estimate: $Q(s, a)$

□ Consider your new sample estimate:

$$y = \text{target}(s') = R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q(s', a')$$

□ Incorporate the new estimate into a running average:

$$Q^{\text{new}}(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \cdot \text{target}(s')$$

Algorithm: (Tabular) Q-Learning

Algorithm:

Start with $Q_0(s, a)$ for all s, a .

Get initial state s

For $k = 1, 2, \dots$ till convergence

 Sample action a , get next state s'

 If s' is terminal:

$\text{target} = R(s, a, s')$

 Sample new initial state s'

 else:

$\text{target} = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$

$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha [\text{target}]$

$s \leftarrow s'$

Learning rate α : $\sum_{t=0}^{\infty} \alpha_t = \infty, \sum_{t=0}^{\infty} \alpha_t^2 < \infty$

Theorem: *Q-learning converges to the optimal action-value function*

$$Q(s, a) \rightarrow Q^*(s, a).$$

How to sample actions?

- Choose random actions?
- Greedy? That is, choose action that maximizes $Q(s, a)$.
- **ϵ -Greedy Policy Exploration:**
 - ◆ With probability ϵ , choose an action at random
 - ◆ With probability $1 - \epsilon$, choose the greedy action

III. DQN

Mnih *et al.*, “Human-level control through deep reinforcement learning,” *Nature*, 2015.

Approximate Q -Learning

- Solve the curse of dimensionality of too big table ($|\mathcal{S}| \times |\mathcal{A}|$) of tabular Q -learning
- Instead of a table, we use a **parametrized** Q -function: $Q_{\theta}(s, a)$ with $\theta = [\theta_1, \dots, \theta_d]^{\top}$.
- Nonlinear neural network approximation of Q -function: Deep Q-Networks (DQN)
 θ : weights of the neural network

Approximate Q -Learning Rule:

□ Remember

$$\text{target}(s') = R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q_{\theta}(s', a')$$

□ Update θ :

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} \left(\frac{1}{2} (Q_{\theta}(s, a) - \text{target}(s'))^2 \right)$$

Why the objective function $\frac{1}{2}(Q_{\theta}(s, a) - \text{target}(s'))^2$ in gradient method?

We can check that the gradient update rule recovers the tabular Q -learning update rule if $Q_{\theta}(s, a) \equiv \theta_{sa}$ for $\theta \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{A}|}$.

Instability of Nonlinear Approximation Using Neural Networks

- ❑ Correlations present in the sequence of observations (episodes)

$$e_t = (s_t, a_t, r_t, s_{t+1})$$

- ❑ Small updates to Q -function may significantly change the policy and therefore change the data distribution.

- ❑ Correlations between $Q(s, a)$ and the target values $\text{target}(s') = R(s, a, s') + \gamma \max_{a' \in \mathcal{A}} Q(s', a')$

Novelty of DQN: stabilize Q-learning for nonlinear approximation with CNN

Experience Replay

- Store agent's experiences $e_t = (s_t, a_t, r_t, s_{t+1})$ at each time-step in a data set $D_t = \{e_1, \dots, e_t\}$ pooled over many episodes into a replay memory. The end of an episode occurs when a terminal state is reached.
- During the inner loop of the algorithm, apply Q-learning updates, or minibatch updates, to samples of experience, $(s, a, r, s') \sim U(D)$, drawn randomly from the pool of stored samples.

Online update \rightarrow off-line

Advantages of experience replay

- ❑ Each step of experience is potentially used in many weight updates, which allows for greater data efficiency.
- ❑ Learning directly from consecutive samples is inefficient, owing to the strong correlations between the samples; randomizing the samples breaks these correlations and therefore reduces the variance of the updates.
- ❑ By using experience replay, the behavior distribution is averaged over many of its previous states, smoothing out learning and avoiding oscillations or divergence in the parameters.
- ❑ When learning by experience replay, it is necessary to learn off-policy (because our current parameters are different to those used to generate the sample), which motivates the choice of Q-learning.

Target Network

Further improving the stability is to use a separate network (target network) for generating the targets $y = r + \gamma \max_{a' \in \mathcal{A}} Q(s', a')$ in the Q-learning update.

Use an older set of weights to compute the targets:

Keeps the target function from changing too quickly.

At iteration i , DQN uses minimizes the following loss function:

$$\min \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a' \in \mathcal{A}} Q_{\theta_i^-}(s', a') - Q_{\theta_i}(s, a) \right)^2 \right]$$

□ θ_i^- : target network parameters at iteration i

□ θ_i : parameters of the Q-network at iteration i

□ The target network parameters θ_i^- are only updated with the Q-network parameters θ_i every C steps and are held fixed between individual updates.

Target Network

- ❑ Every C updates we clone the network Q to obtain a target network \hat{Q} and use \hat{Q} for generating the Q-learning targets y for the following C updates to Q .
- ❑ This modification makes the algorithm more stable compared to standard online Q-learning.
- ❑ Generating the targets using an older set of parameters adds a delay between the time an update to Q is made and the time the update affects the targets y , making divergence or oscillations much more unlikely.

Other Details of DQN

- Downsampling: 210×160 game images to ones of 84×84 ($s_t \rightarrow \phi(s_t)$).
- CNN: 5 layers: 3 convolution layers + 2 full connection layers

Layer	Input	Filter size	Stride	Num filters	Activation	Output
conv1	84x84x4	8x8	4	32	ReLU	20x20x32
conv2	20x20x32	4x4	2	64	ReLU	9x9x64
conv3	9x9x64	3x3	1	64	ReLU	7x7x64
fc4	7x7x64			512	ReLU	512
fc5	512			18	Linear	18

- Uses RMSProp instead of vanilla SGD
Optimization in RL really matters.
- It helps to anneal the exploration rate:
Start ϵ at 1 and anneal it to 0.1 or 0.05 over the first million frames.

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ε select a random action a_t

otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For

End For

Thank you!